

Stat 5200: Analysis of Designed Experiments

Jesse Wheeler

Overview

Objectives and overview

- This material is based Chapter 3 on the textbook “A First Course in Design and Analysis of Experiments” (Oehlert, 2000).
- Chapter outcomes include:
 - TODO

Completely Randomized Designs

Completely Randomized Design

- Completely randomized design (CRD) is the simplest design for comparing treatments and is sometimes the preferred design.
- Setup:
 - There are g treatment groups.
 - N experimental units (EU)s available.
 - You want sample sizes of $n_1, n_2, \dots, n_g$, with $\sum_{i=1}^g n_i = N$.

Completely Randomized Design II

- **Idea:** All possible sub-group composition are equally as likely.

Performing a Completely Randomized Design

- Randomly choose n_1 EU's for treatment 1.
 - Randomly choose n_2 EU's for treatment 2.
 - ...
 - Randomly choose n_{g-1} EU's for treatment $g - 1$.
 - The remaining n_g EU's are assigned to treatment g .
-
- On a computer, this can be done by permuting $1, 2, \dots, N$, and assigning the first n_1 numbers to treatment 1, the next n_2 to treatment 2, etc.

Advantages and Disadvantages

- Advantages
 - Simple to do; all you need is one random permutation.
 - Simple to explain and understand.
 - Simple to analyze.
 - No restrictions on the sample sizes (other designs may restrict n_i).
 - Easy with missing values.
- Disadvantages
 - Not always possible or cost-effective.
 - The accuracy of the experiment depends on variability among the EUs, not just the variability among treatment groups. Thus, it may require more EUs (\$\$\$\$).

Example: Acid Rain

Wood and Bormann (1974) studied the effect of acid rain on trees. One experiment used 240 six-week-old yellow birch seedlings. These seedlings were divided into five groups of 48 *at random*, and the seedlings within each group received an acid mist treatment 6 hours a week for 17 weeks. The five treatments differed by mist pH: 4.7, 4.0, 3.3, 3.0, and 2.3. Otherwise, the seedlings were treated identically. After the 17 weeks, the seedlings were weighed, and total plant (dry) weight was taken as response.

Questions:

- What type of design is this, and can you describe why?
- What are the EUs, and how many are there?
- What was the response?

Example: Acid Rain II

- What compromises may have been made in this case?

Like many experiments, the researchers necessarily have to narrow down an enormous question to a workable experiment using artificial acid rain under controlled conditions. For any experiment, a considerable amount of non-statistical background work goes into the planning even the simplest experiment.

Models and Parameters

Introduction and definitions

- All formal statistical analysis is based on probability models.
- *EX*: The number of heads in ten tosses of a fair coin would be $\text{Binomial}(10, 0.5)$. Here, the model depends on 2 parameters: $N = 10$ the number of trials, $p = 0.5$ is the probability of heads.
- In an experiment, we may posit several different models for the data, all with *unknown* parameters. One of our objectives is often to describe which model is the best description of the data, and make *inference* about parameters.

Definition: Inference

Using data to learn properties of a probability distribution. In our case, using data to estimate parameters of a model.

Introduction and definitions II

- The models we will consider for experimental design will have two basic parts:
 - A description of the expected value, or mean, of the data. Sometimes called “model for the means”.
 - A description of how data vary around the treatment means. Sometimes called ‘model for the errors’.
- For the errors, we will generally assume that they are independent for different data values, have mean zero, and all have a common variance σ^2 .
- To perform formal tests and obtain confidence intervals, we will further assume a distribution of the errors. Generally, we assume that the errors follow a **Gaussian** (Normal) distribution.

Introduction and definitions III

- For a CRD, we start with two basic models.
 - Recall: g treatment groups, N EUs.
1. **The Full Model** (separate group means)

Introduction and definitions IV

2. The Reduced Model (single mean model)

Defining overall mean

- It's sometimes useful to express the group means μ_i as:

$$\mu_i = \mu^* + \alpha_i.$$

In this case, we call μ^* the *overall mean*, and the α_i 's are the *i*th *treatment effects*.

- This introduces an extra parameter ($g+1$ total), which we need to address. Generally, we do this by defining μ^* and require that $\alpha_i = \mu_i - \mu^*$.

Defining overall mean II

- One classic choice is to define μ^* as the mean of the treatment means:

Defining overall mean III

- An alternative choice (that can sometimes make hand-work simpler) instead describes μ^* as the weighted average:

Defining overall mean IV

- Another common choice in software is to set $\mu^* = \mu_j$, for some $j \in 1 : g$. Many software just use the last mean: $\mu^* = \mu_g$.

Model Recap

However you choose to define μ^* and α_i 's, we end up with the following models for the data:

- *The full model:*

$$y_{ij} = \mu^* + \alpha_i + \epsilon_{ij}$$

- *The reduced model:*

$$y_{ij} = \mu + \epsilon_{ij}$$

Parameter Estimation

Estimating parameters

- We want to use data to estimate μ_i 's and the variance σ^2 . From these, we'll compute μ^* and α_i 's.
 - We will focus on computing **unbiased** estimates of these parameters.
 - Estimate for μ_i :
-
- The estimate for μ^* is fixed, so it depends on how we defined it.

Estimating parameters II

- Unweighted average
- Weighted average
- Reference Group

Estimating parameters III

- The only parameter left to estimate is σ^2 :

Estimating parameters IV

- Above, we have described **point estimates**. In some sense, these are our “best guess” as to the value of a particular parameter.
- We also want some measure of uncertainty of this estimate. The standard frequentist way of doing this is a **confidence interval**.
- Confidence intervals for α are not very useful, because it depends on how we define them. Confidence intervals for group means are typically defined as:

Comparing Models: ANOVA

Analysis of Variance

- In our analysis of a CRD, our models focus on the *mean responses* of the treatment groups.
- A basic question we might want answered: **Are all the means the same, or are some different?**
- We can express this same question mathematically:

Analysis of Variance II

- The basic idea for many statistical tests (including ANOVA) follows *Occam's Razor*: we use the simplest model that is consistent with the data; we will only use a more complicated model if the data indicate that the more complicated model is needed.
- How do we know if a more complex model is needed? We compare **residuals**.

Analysis of Variance III

- The lower the SSR, the more of the variance in the data is explained by the model.
- Importantly, the SSR for the reduced (single mean) model can *never* be lower than the separate model, and in practice, is nearly always larger.
- What we need now is a way to consistently decide when the separate means model is sufficiently “better” than the single means model.
- To do so, we will set up a statistical hypothesis test:

ANOVA mechanics

- ANOVA works by partitioning the total variance in the data into parts that mimic the model.
- For instance, in the separate mean model, the model says the data differ from the grand mean by a random error and a treatment effect:

ANOVA will have a similar decomposition: one part dealing with group means, one part with deviations from those means:

- Recall from Eq. ??, we have:

$$y_{ij} - \bar{y}_{\bullet\bullet} = \hat{\alpha}_i + r_{ij},$$

ANOVA mechanics III

- Finally, we can connect this back to the residuals.
- The Sum of Squared Residuals (SSR) depended on our model:

$$SSR_{\text{single means}} = \sum_{i=1}^g \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{\bullet\bullet})^2 = SS_T.$$

$$SSR_{\text{separate means}} = \sum_{i=1}^g \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i\bullet})^2 = SS_E$$

Calibrating Effect Size

- We have established:
 - We will compare single mean against separate mean by comparing their SSRs.
 - Comparing SSR can be done by decomposing total error into two parts: $SS_T = SS_{T_{rt}} + SS_E$, and if the separate model is “much better” than the $SS_{T_{rt}}$ will be large.
- **How large is large enough?**

Calibrating Effect Size II

- The answer also lies in the decomposition of sum of squares.

ANOVA variable recap

- Sum of Square Residuals:

$$SSR = \sum_{i=1}^g \sum_{j=1}^{n_i} (y_{ij} - \hat{y}_{ij})^2$$

- Treatment Sum of Squares:

$$SS_{Trt} = \sum_{i=1}^g \sum_{j=1}^{n_i} (\bar{y}_{i\bullet} - \bar{y}_{\bullet\bullet})^2 = \sum_{i=1}^g n_i \hat{\alpha}_i^2.$$

- Error Sum of Squares:

$$SS_E = \sum_{i=1}^g \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{i\bullet})^2$$

ANOVA variable recap II

- Total Sum of Squares:

$$SS_T = SS_{Trt} + S_E = \sum_{i=1}^g \sum_{j=1}^{n_i} (y_{ij} - \bar{y}_{\bullet\bullet})^2$$

- F-statistic:

$$\frac{SS_{Trt}/(g-1)}{SS_E/(N-g)} \sim F_{g-1, N-g}.$$

References and Acknowledgements

Oehlert GW (2000). *A First Course in Design and Analysis of Experiments*. W. H. Freeman. ISBN 0-7167-3510-5.

Wood T, Bormann F (1974). “The effects of an artificial acid mist upon the growth of *Betula alleghaniensis* Britt.” *Environmental Pollution* (1970), **7**(4), 259–268.

- Compiled on September 1, 2026 using R version 4.5.0.
- Licensed under the [Creative Commons Attribution-NonCommercial](#) license. Please share and remix non-commercially, mentioning its origin.



References and Acknowledgements II

- We acknowledge [students and instructors for previous versions of this course / slides](#).