

Stat 5200, Module 4: Factorial Design and Analysis

Jesse Wheeler

Overview

Objectives and overview

- This material is based on Chapters 8–10 of the textbook “A First Course in Design and Analysis of Experiments” (Oehlert, 2000).
- Module outcomes include:
 - You will understand complete and incomplete factorial treatment structures.
 - You will develop an intuitive and formal understanding of interaction effects.
 - You will be able to specify and fit additive and non-additive two-way ANOVA models.
 - You will learn how to construct and interpret main effect contrasts and interaction contrasts.
 - You will know how to fit and analyze factorial designs in R using `aov()` and `emmeans()`.

Factorial Treatment Structure

Introduction to Factorial Structures

- In many experiments, we want to evaluate the effect of two or more explanatory factors simultaneously.
- A **factorial treatment structure** defines treatments as combinations of levels of each factor.
- Suppose Factor A has a levels and Factor B has b levels.
- A **complete factorial design** includes all $g = a \times b$ possible treatment combinations.

Introduction to Factorial Structures II

Example 1: 2×3 Diet Experiment

- Factor 1: Protein Source (Beef, Cereal, Pork) $\implies b = 3$.
- Factor 2: Protein Intake (High, Low) $\implies a = 2$.
- Total treatment combinations: $g = 2 \times 3 = 6$ unique diets.

Introduction to Factorial Structures III

Example 2: 2×2 Fertilizer Experiment

- Factor A: Superphosphate (P) $\rightarrow$ Levels: {Not Added, Added}
- Factor B: Potash (K) $\rightarrow$ Levels: {Not Added, Added}
- $g = 2 \times 2 = 4$ treatments:
 1. None (P_1K_1)
 2. Superphosphate only (P_2K_1)
 3. Potash only (P_1K_2)
 4. Both Superphosphate and Potash (P_2K_2)

Main Effects and Interactions

- We are generally interested in both the *main effects* of each treatment, but also the *interactions* between treatments. Consider the example below (Example 2):

	Beef	Cereal	Pork	Marginal Row Mean
High Intake	$\mu_{B,H}$	$\mu_{C,H}$	$\mu_{P,H}$	μ_H
Low Intake	$\mu_{B,L}$	$\mu_{C,L}$	$\mu_{P,L}$	μ_L
Marginal Col Mean	μ_B	$\bar{\mu}_C$	μ_L	μ (Grand Mean)

Main Effects and Interactions II

- **Main Effects:** Evaluated by comparing the marginal row means ($\bar{\mu}_{1\bullet}$ vs. $\bar{\mu}_{2\bullet}$) or the marginal column means ($\bar{\mu}_{\bullet 1}$, $\bar{\mu}_{\bullet 2}$, $\bar{\mu}_{\bullet 3}$). For example, what is the difference in response between high and low intake treatments?
- **Interactions:** Evaluated by looking at differences between cells. For example, is the difference between High and Low intake the same for Beef as it is for Cereal?

$$(\mu_{B,H} - \mu_{B,L}) \stackrel{?}{=} (\mu_{C,H} - \mu_{C,L})$$

Main Effects and Interactions III

- While we *could* formalize and test all these specific comparisons using individual contrasts, this can be tedious and less interpretable in some cases.
- Instead, we will introduce the **Two-Way ANOVA** as an efficient way to test for overall main effects and interactions simultaneously, with highly interpretable results.
- Like one-way ANOVA, we will express the total sum of squares in terms of treatments and errors.

Notation for Two-Way Factorials

- To identify observations in a two-way design, we use three subscripts:

$$y_{ijk}$$

- $i = 1, \dots, a$: level of Factor A
- $j = 1, \dots, b$: level of Factor B
- $k = 1, \dots, n$: replication number within treatment combination (i, j)
- Summary Notation:
 - Sample mean for treatment (i, j) : $\bar{y}_{ij\bullet} = \frac{1}{n} \sum_{k=1}^n y_{ijk}$
 - Marginal mean for level i of Factor A:
$$\bar{y}_{i\bullet\bullet} = \frac{1}{bn} \sum_{j=1}^b \sum_{k=1}^n y_{ijk}$$

Notation for Two-Way Factorials II

- Marginal mean for level j of Factor B:

$$\bar{y}_{\bullet j \bullet} = \frac{1}{an} \sum_{i=1}^a \sum_{k=1}^n y_{ijk}$$

- Overall grand mean: $\bar{y}_{\bullet \bullet \bullet} = \frac{1}{abn} \sum_{i=1}^a \sum_{j=1}^b \sum_{k=1}^n y_{ijk}$

The Two-Way ANOVA Model

Model Specification

- For completely randomized designs, we will start with **balanced factorial** treatments. All $n_{ij} = n$ for all i, j .
- We model the response y_{ijk} as:

$$y_{ijk} = \mu + \alpha_i + \beta_j + (\alpha\beta)_{ij} + \epsilon_{ijk}$$

where:

- μ : Overall grand mean.
- α_i : Main effect of level i of Factor A.
- β_j : Main effect of level j of Factor B.
- $(\alpha\beta)_{ij}$: Interaction effect between level i of A and level j of B.
- $\epsilon_{ijk} \stackrel{iid}{\sim} N(0, \sigma^2)$: Random errors.

Model Specification II

- Parameter Identification Restrictions (Sum-to-Zero Constraints):

$$\sum_{i=1}^a \alpha_i = \sum_{j=1}^b \beta_j = \sum_{i=1}^a (\alpha\beta)_{ik} = \sum_{j=1}^b (\alpha\beta)_{kj} = 0 \text{ for all } k$$

- Hierarchy of Models:

- Full (Non-additive) Model:**

$$y_{ijk} = \mu + \alpha_i + \beta_j + (\alpha\beta)_{ij} + \epsilon_{ijk}$$

- Additive Model:** $y_{ijk} = \mu + \alpha_i + \beta_j + \epsilon_{ijk}$

- Main Effect Models:** $y_{ijk} = \mu + \alpha_i + \epsilon_{ijk},$

$$y_{ijk} = \mu + \beta_j + \epsilon_{ijk}$$

- Constant Model:** $y_{ijk} = \mu + \epsilon_{ijk}$

Degrees of Freedom

- The degrees of freedom for each main effect is computed as we would in a one-way ANOVA:
 - Main effect A : $(a - 1)$
 - Main effect B : $(b - 1)$
 - Interaction effect (AB) : $(a - 1)(b - 1)$.

Understanding Interaction Effects

What is an Interaction Effect?

- **Definition:** An interaction effect occurs when the effect of one factor on the response depends on the level of another factor.
- Ask yourself: *"Does the difference between levels of Factor A stay the same regardless of Factor B?"*
 - If **YES**: The factors are **additive** (No Interaction).
 - If **NO**: There is an **interaction** between factors.

What is an Interaction Effect? II

Intuition: "Difference of Differences"

Consider a 2×2 design with population means μ_{ij} :

- Simple main effect of A when Factor B is at level 1: $\mu_{11} - \mu_{21}$
- Simple main effect of A when Factor B is at level 2: $\mu_{12} - \mu_{22}$

The **interaction effect** is the difference between these two simple effects:



$$(\mu_{11} - \mu_{21}) - (\mu_{12} - \mu_{22})$$

If this difference equals zero, the effect of A is constant across levels of B. If non-zero, an interaction exists.

Graphical Diagnosis: Interaction Plots

- An **interaction plot** displays treatment group means across levels of one factor, with lines connecting levels of the second factor.
- **Parallel lines** $\implies$ No Interaction (Additive).
- **Non-parallel or crossing lines** $\implies$ Interaction present.

Graphical Diagnosis: Interaction Plots II



ANOVA Mechanics and Hypothesis Testing

Two-Way ANOVA Table (Balanced Design)

- For a balanced design (n observations per cell, $N = abn$ total):

Source	DF	SS	MS	F-Statistic
Factor A	$a - 1$	SS_A	MS_A	$F_A = \frac{MS_A}{MS_E}$
Factor B	$b - 1$	SS_B	MS_B	$F_B = \frac{MS_B}{MS_E}$
Interaction	$(a - 1)(b - 1)$	SS_{AB}	MS_{AB}	$F_{AB} = \frac{MS_{AB}}{MS_E}$
Error	$ab(n - 1)$	SS_E	MS_E	
Total	$abn - 1$	SS_T		

See compiled notes for detailed formulas.

Two-Way ANOVA Table (Balanced Design) II

Sequence of Testing: Top-Down Rule

1. **Step 1: Test Interaction ($H_0 : (\alpha\beta)_{ij} = 0$ for all i, j)**
 - If **Significant**: Stop! The model is non-additive. Do not interpret main effect F -tests in isolation. Examine simple main effects or interaction contrasts.
2. **Step 2: Test Main Effects**
 - If interaction is non-significant, proceed to test $H_0 : \alpha_i = 0$ and $H_0 : \beta_j = 0$ independently using the main effect F -tests.

Applied Examples in R

Example 1: Significant Interaction (ToothGrowth)

- The built-in `ToothGrowth` dataset measures tooth length (`len`) in guinea pigs under two factors:
 - **Supplement** (`supp`): Orange Juice (OJ) or Ascorbic Acid (VC).
 - **Dose** (`dose`): 0.5, 1.0, or 2.0 mg/day.
- First, we must convert dose from a numeric variable into a factor so R treats it as categorical.

```
# Load data and convert dose to a factor
```

```
data(ToothGrowth)
```

```
ToothGrowth$dose <- as.factor(ToothGrowth$dose)
```


Example 1: Significant Interaction (ToothGrowth) II

Visualizing the Interaction

We start with an interaction plot. Crossing or converging lines suggest an interaction effect.

```
interaction.plot(  
  x.factor = ToothGrowth$dose,  
  trace.factor = ToothGrowth$supp,  
  response = ToothGrowth$len,  
  type = "b", pch = c(16, 17),  
  col = c("orange", "blue"),  
  xlab = "Dose (mg/day)",  
  ylab = "Mean Tooth Length",  
  trace.label = "Supplement"  
)
```

Example 1: Significant Interaction (ToothGrowth) III



Example 1: Significant Interaction (ToothGrowth) IV

Fitting the Model and Testing (Top-Down Rule)

We fit the full model with the * operator and check the ANOVA table.

```
mod_tooth <- aov(len ~ supp * dose, data = ToothGrowth)
summary(mod_tooth)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
supp	1	205.4	205.4	15.572	0.000231	***
dose	2	2426.4	1213.2	92.000	< 2e-16	***
supp:dose	2	108.3	54.2	4.107	0.021860	*
Residuals	54	712.1	13.2			

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Example 1: Significant Interaction (ToothGrowth) V

- **Interpretation:** The `supp:dose` interaction is significant ($p = 0.02$). Following our Top-Down Rule, we **stop** and do not interpret the main effects blindly. The effect of the supplement depends entirely on the dose.

Example 1: Significant Interaction (ToothGrowth) VI

Analyzing Simple Effects with `emmeans`

Because the interaction is significant, we should look at the "simple effects"—the effect of the supplement *at each specific level* of the dose.

Example 1: Significant Interaction (ToothGrowth) VII

```
library(emmeans)
# The "/" operator tells emmeans to compare 'supp' AT each level
emm_tooth <- emmeans(mod_tooth, pairwise ~ supp | dose)

# View the pairwise comparisons
emm_tooth$contrasts
```

dose = 0.5:

contrast	estimate	SE	df	t.ratio	p.value
OJ - VC	5.25	1.62	54	3.233	0.0021

dose = 1:

contrast	estimate	SE	df	t.ratio	p.value
OJ - VC	5.93	1.62	54	3.651	0.0006

dose = 2:

contrast	estimate	SE	df	t.ratio	p.value
----------	----------	----	----	---------	---------

Demystifying the `emmeans` Syntax

- In a one-way ANOVA, we could use `emmeans(model, ~ Factor)` to get the estimated means for each group.
- In a two-way ANOVA with a significant interaction, we must examine **simple main effects**. We need to know the effect of one factor *at specific levels* of the other factor.

Breaking down: `pairwise ~ supp | dose`

- `pairwise ~ supp`: This tells R to calculate the pairwise differences between the levels of the Supplement factor (OJ vs. VC).
- `| (The Pipe/Bar)`: In R formulas, the vertical bar means "given" or "conditioned on."
- `dose`: This is the grouping variable. It tells R to split the analysis up by the levels of Dose.

Example 2: No Interaction (Additive Model)

- The built-in npk dataset records pea yields based on applying Nitrogen (N) and Phosphate (P).
- Let's evaluate if Nitrogen and Phosphate interact to affect yield.

```
NP <- read.csv('NP.csv')  
str(NP)
```

```
'data.frame': 24 obs. of 3 variables:  
 $ yield: num 49.5 62.8 46.8 57 59.8 58.5 55.5 56 62..  
 $ P : int 1 1 0 0 0 1 0 1 1 1 ...  
 $ N : int 0 1 0 1 1 1 0 0 0 1 ...
```

Example 2: No Interaction (Additive Model) II

```
mod_NP <- aov(yield ~ N * P, data = NP)
summary(mod_NP)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
N	1	189.3	189.28	5.758	0.0263	*
P	1	8.4	8.40	0.256	0.6187	
N:P	1	21.3	21.28	0.647	0.4305	
Residuals	20	657.4	32.87			

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Example 2: No Interaction (Additive Model) III

- **Interpretation:** The N:P interaction is not significant ($P = 0.4305$). That is, there is not strong evidence in the data that the effect of Nitrogen inclusions on Mean Yield depends on the inclusion of Phosphate.
- We can examine this further with the interaction plot

Example 2: No Interaction (Additive Model) IV

```
# Visualize the data
interaction.plot(
  x.factor = NP$N,
  trace.factor = NP$P,
  response = NP$yield,
  type = "b", pch = c(16, 17),
  col = c("darkgreen", "brown"),
  xlab = "Nitrogen (N)", ylab = "Mean Yield",
  trace.label = "Phosphate (P)"
)
```

Example 2: No Interaction (Additive Model) V



Example 2: No Interaction (Additive Model) VI

- Not significant interaction. Now what?
 - Historically, we might want to re-fit an **additive model** to conserve degrees of freedom.
Problem: This is called *post-selection inference*, and can mean we have inflated type-I error rates.
- Current practice: we are more concerned about the increased Type-I error rate than we are of improving statistical power. We will interpret P-values in the context of accounting for interactions, even if they are non-significant.

Example 2: No Interaction (Additive Model) VII

(Historic approach): Fitting the Additive Model

We use the + operator to fit a model assuming the effects of N and P are strictly independent.

Example 2: No Interaction (Additive Model) VIII

```
# Fit additive model (no interaction)
```

```
mod_NP_add <- aov(yield ~ N + P, data = NP)
```

```
summary(mod_NP_add)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
N	1	189.3	189.28	5.857	0.0247	*
P	1	8.4	8.40	0.260	0.6155	
Residuals	21	678.7	32.32			

```
---
```

```
Signif. codes:
```

```
0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Example 2: No Interaction (Additive Model) IX

- **Conclusion:** Nitrogen (N) has a significant main effect on yield ($p = 0.0247$), while Phosphate (P) does not ($p = 0.6155$). Because the model is additive, we conclude Nitrogen increases yield regardless of whether Phosphate is applied or not.
- We could continue to extract treatment group means and treatment effects using the `emmeans` package as needed (see below).

Extracting Treatment Means and Effects

- Following modern practice, we retain the full model (`mod_NP`) even though the interaction was not significant.
- We can use `emmeans` to extract the estimated marginal means (treatment means) for Nitrogen, averaging over the levels of Phosphate.

Extracting Treatment Means and Effects II

```
library(emmeans)

# Extract Estimated Marginal Means for Nitrogen (N)
# mod_NP is FULL model, not additive
emm_N <- emmeans(mod_NP, ~ N)
emm_N
```

N	emmean	SE	df	lower.CL	upper.CL
0	52.1	1.66	20	48.6	55.5
1	57.7	1.66	20	54.2	61.1

Results are averaged over the levels of: P
Confidence level used: 0.95

Extracting Treatment Means and Effects III

Calculating the Treatment Effect

To find the actual size of the treatment effect (the difference between adding Nitrogen vs. not adding Nitrogen), we pass our means into the `pairs()` function.

```
# Test the pairwise difference (Treatment Effect)
```

```
pairs(emm_N)
```

contrast	estimate	SE	df	t.ratio	p.value
N0 - N1	-5.62	2.34	20	-2.400	0.0263

Results are averaged over the levels of: P

Extracting Treatment Means and Effects IV

- **Interpretation:** The estimated mean yield without Nitrogen is 52.1, and with Nitrogen is 57.7. The `pairs()` output shows that applying Nitrogen increases the mean yield by an estimated 5.62 units, and this main effect is statistically significant.

Factorial Contrasts

Main Effects vs. Interaction Contrasts

- A contrast in a factorial design is a linear combination of treatment means $\sum_i \sum_j w_{ij} \mu_{ij}$ subject to $\sum_i \sum_j w_{ij} = 0$.
- **Main Effect Contrasts for A:** Weights for a given level of A are identical across all levels of B.

$$\sum_{i=1}^a w_i = 0, \quad w_{ij} = w_i \quad \text{for all } j$$

Main Effects vs. Interaction Contrasts II

- **Interaction Contrasts ("Contrast of Contrasts")**: Weights are constructed as products of main effect contrast weights:
 $w_{ij} = w_{A,i} \times w_{B,j}$. Row sums and column sums of weights must **all equal zero**:

$$\sum_{i=1}^a w_{ij} = 0 \text{ for all } j, \quad \text{and} \quad \sum_{j=1}^b w_{ij} = 0 \text{ for all } i$$

Main Effects vs. Interaction Contrasts III

Example: 2×2 Factorial Weight Matrices

Main Effect A

(W_A)

1	1
-1	-1

Main Effect B

(W_B)

1	-1
1	-1

Interaction (W_{AB})

1	-1
-1	1

Advanced Topics in Practice (R)

Unbalanced Factorials & Simple Main Effects

Types of Sums of Squares (3-Way Factorial)

- When cell sizes are unequal ($n_{ijk} \neq n$), factors lose orthogonality.
- The choice of **Sums of Squares (SS)** dictates which model terms are held constant in the reduced vs. full model comparison ($SS_{\text{Effect}} = SS_{\text{Error, reduced}} - SS_{\text{Error, full}}$).

Effect	Type I (Sequential)	Type II (Hierarchical)	Type III
<i>A</i>	$SS(A)$	$SS(A \mid B, C)$	$SS(A \mid B, C)$
<i>B</i>	$SS(B \mid A)$	$SS(B \mid A, C)$	$SS(B \mid A, C)$
<i>C</i>	$SS(C \mid A, B)$	$SS(C \mid A, B)$	$SS(C \mid A, B)$
<i>AB</i>	$SS(AB \mid A, B, C)$	$SS(AB \mid A, B, C, AC, BC)$	$SS(AB \mid A, B, C, AC, BC)$
<i>AC</i>	$SS(AC \mid A, B, C, AB)$	$SS(AC \mid A, B, C, AB, BC)$	$SS(AC \mid A, B, C, AB, BC)$
<i>BC</i>	$SS(BC \mid A, B, C, AB, AC)$	$SS(BC \mid A, B, C, AB, AC)$	$SS(BC \mid A, B, C, AB, AC)$
<i>ABC</i>	$SS(ABC \mid A, B, C, AB, AC, BC)$	$SS(ABC \mid A, B, C, AB, AC, BC)$	$SS(ABC \mid A, B, C, AB, AC, BC)$

Types of Sums of Squares (3-Way Factorial) II

Table 1: Comparison of Sums of Squares conditional notation for factors A, B, and C.

Types of Sums of Squares (3-Way Factorial) III

Which Type Should You Use in Practice?

- **Type I:** Avoid unless testing a strict hierarchical sequence (e.g., covariate first). Order of terms in `aov()` changes the output!
- **Type II:** Best default when you are confident higher-order interactions are zero ($AB = 0$). Maximizes statistical power.
- **Type III:** Standard for factorial experiments. Tests equal weighting across all treatment cells regardless of sample size differences.

Unbalanced Data: Practical Workflow in R

- Base R's `summary(aov())` defaults to Type I SS, making it dangerous for unbalanced data.
- **Solution:** Use `joint_tests()` from `emmeans`. It provides Type III hypothesis tests safely without requiring you to change R's global options.

```
library(emmeans)

# Fit standard model on unbalanced data
mod_unbal <- aov(yield ~ P * K, data = unbalanced_data)

# 1. Run Type III ANOVA tests safely
joint_tests(mod_unbal)
```

Unbalanced Data: Practical Workflow in R II

Practical Decision Rules

1. Look at the Interaction p-value first:

- If $p \geq 0.05$ (No interaction): Interpret overall main effects and run overall pairwise comparisons: `emmeans(mod, P)`.
- If $p < 0.05$ (Interaction present): **Stop!** Do not interpret global main effects. You must analyze **Simple Main Effects**.

Analyzing Simple Main Effects

- **Definition:** A **simple main effect** tests the effect of Factor A at **one specific fixed level** of Factor B.
- When an interaction exists, stating "Superphosphate increases yield" is incomplete. You must state "Superphosphate increases yield **when Potash is absent**, but has no effect **when Potash is present**."
- Use the `by` argument in `emmeans` to slice the analysis.

Analyzing Simple Main Effects II

```
# Obtain EMMs for the interaction
```

```
emm_int <- emmeans(mod_unbal, ~ P * K)
```

```
# Test the simple main effect of Potash (K) WITHIN each level of
```

```
contrast(emm_int, method = "pairwise", by = "P")
```

```
# Plot simple effects visually
```

```
emmip(mod_unbal, K ~ P)
```

Multiple Comparisons: Beyond the F-Test

Multiple Comparisons in Practice

- A significant ANOVA interaction tells you *some* difference exists, but testing all pairs blindly increases your **family-wise error rate** (false positive rate).
- The `emmeans` package provides a streamlined pipeline for post-hoc testing with appropriate data-dredging penalties.

```
# Assuming fert_mod is our fitted two-way model  
emm_int <- emmeans(fert_mod, ~ P * K)  
  
# Tukey adjustment for all pairwise comparisons  
pairs(emm_int, adjust = "tukey")
```

Multiple Comparisons in Practice II

Visualizing Contrasts

- Staring at massive tables of p-values makes it difficult to assess effect sizes.
- You can pass your contrast object directly into a plot to visually see which adjusted confidence intervals exclude zero.

```
# Generates a forest plot of the differences between groups  
plot(pairs(emm_int, adjust = "tukey")) +  
  ggplot2::geom_vline(xintercept = 0, linetype = "dashed", color = "red") +  
  ggplot2::labs(title = "Tukey-Adjusted Pairwise Differences")
```

Resource Constraints: Unreplicated & Fractional Designs

Unreplicated Factorials ($n = 1$)

- Experiments are expensive. Sometimes we can only afford a single replicate per treatment combination ($n = 1$).
- **The Problem:** Fitting a full factorial model uses up all available degrees of freedom. The error degrees of freedom (df_E) becomes 0, so R returns NA for all standard errors and p-values.

```
# Fits main effects, 2-way, 3-way, and 4-way interactions  
mod_unrep <- aov(yield ~ A * B * C * D, data = unrep_data)  
  
# F-values and p-values will return as NA  
summary(mod_unrep)
```

Unreplicated Factorials ($n = 1$) II

The Fix: Pooling into Error

- Relying on the **Sparsity of Effects Principle**, we assume high-order interactions (e.g., 3-way and 4-way) are effectively zero (just random noise).
- We restrict the model, pushing those negligible higher-order terms into the residual error to mathematically generate degrees of freedom.

Unreplicated Factorials ($n = 1$) III

```
# Restrict the model to main effects and 2-way interactions only  
# The "^2" syntax drops 3-way and 4-way terms into the error  
mod_pooled <- aov(yield ~ (A + B + C + D)^2, data = unrep_data)  
  
# Now R can calculate p-values for main and 2-way effects!  
summary(mod_pooled)
```

Fractional Factorials & Confounding

- A **fractional factorial** design runs only a structured fraction of the full combinations (e.g., 16 runs instead of 32 for a 5-factor design).
- **How R handles it:** R automatically recognizes the missing runs and returns NA for the interaction terms it cannot mathematically estimate.
- **Aliasing:** Because we skipped runs, certain effects are perfectly correlated (confounded).
 - *Intuition:* If we design a half-fraction by defining Factor C's levels mathematically as $A \times B$, R cannot tell if an observed change in yield is caused by the main effect of C, or the interaction of A and B.

Fractional Factorials & Confounding II

```
# Generate a mock half-fraction dataset ( $2^{(3-1)}$ )
library(FrF2)
frac_data <- FrF2(nruns = 4, nfactors = 3, factor.names = c("A",
frac_data$yield <- c(14, 18, 22, 19)

frac_mod <- aov(yield ~ A * B * C, data = frac_data)

# This outputs the exact matrix showing which main effects
# are mathematically identical to which interaction terms:
alias(frac_mod)

Model :
yield ~ A * B * C

Complete :
              (Intercept) A1 B1 C1
A1:B1      0             0  0  1
```

References and Acknowledgements

Oehlert GW (2000). *A First Course in Design and Analysis of Experiments*. W. H. Freeman. ISBN 0-7167-3510-5.

- Compiled on October 8, 2026 using R version 4.6.1.
- Licensed under the [Creative Commons Attribution-NonCommercial](#) license. Please share and remix non-commercially, mentioning its origin.
- We acknowledge [students and instructors for previous versions of this course / slides](#).

